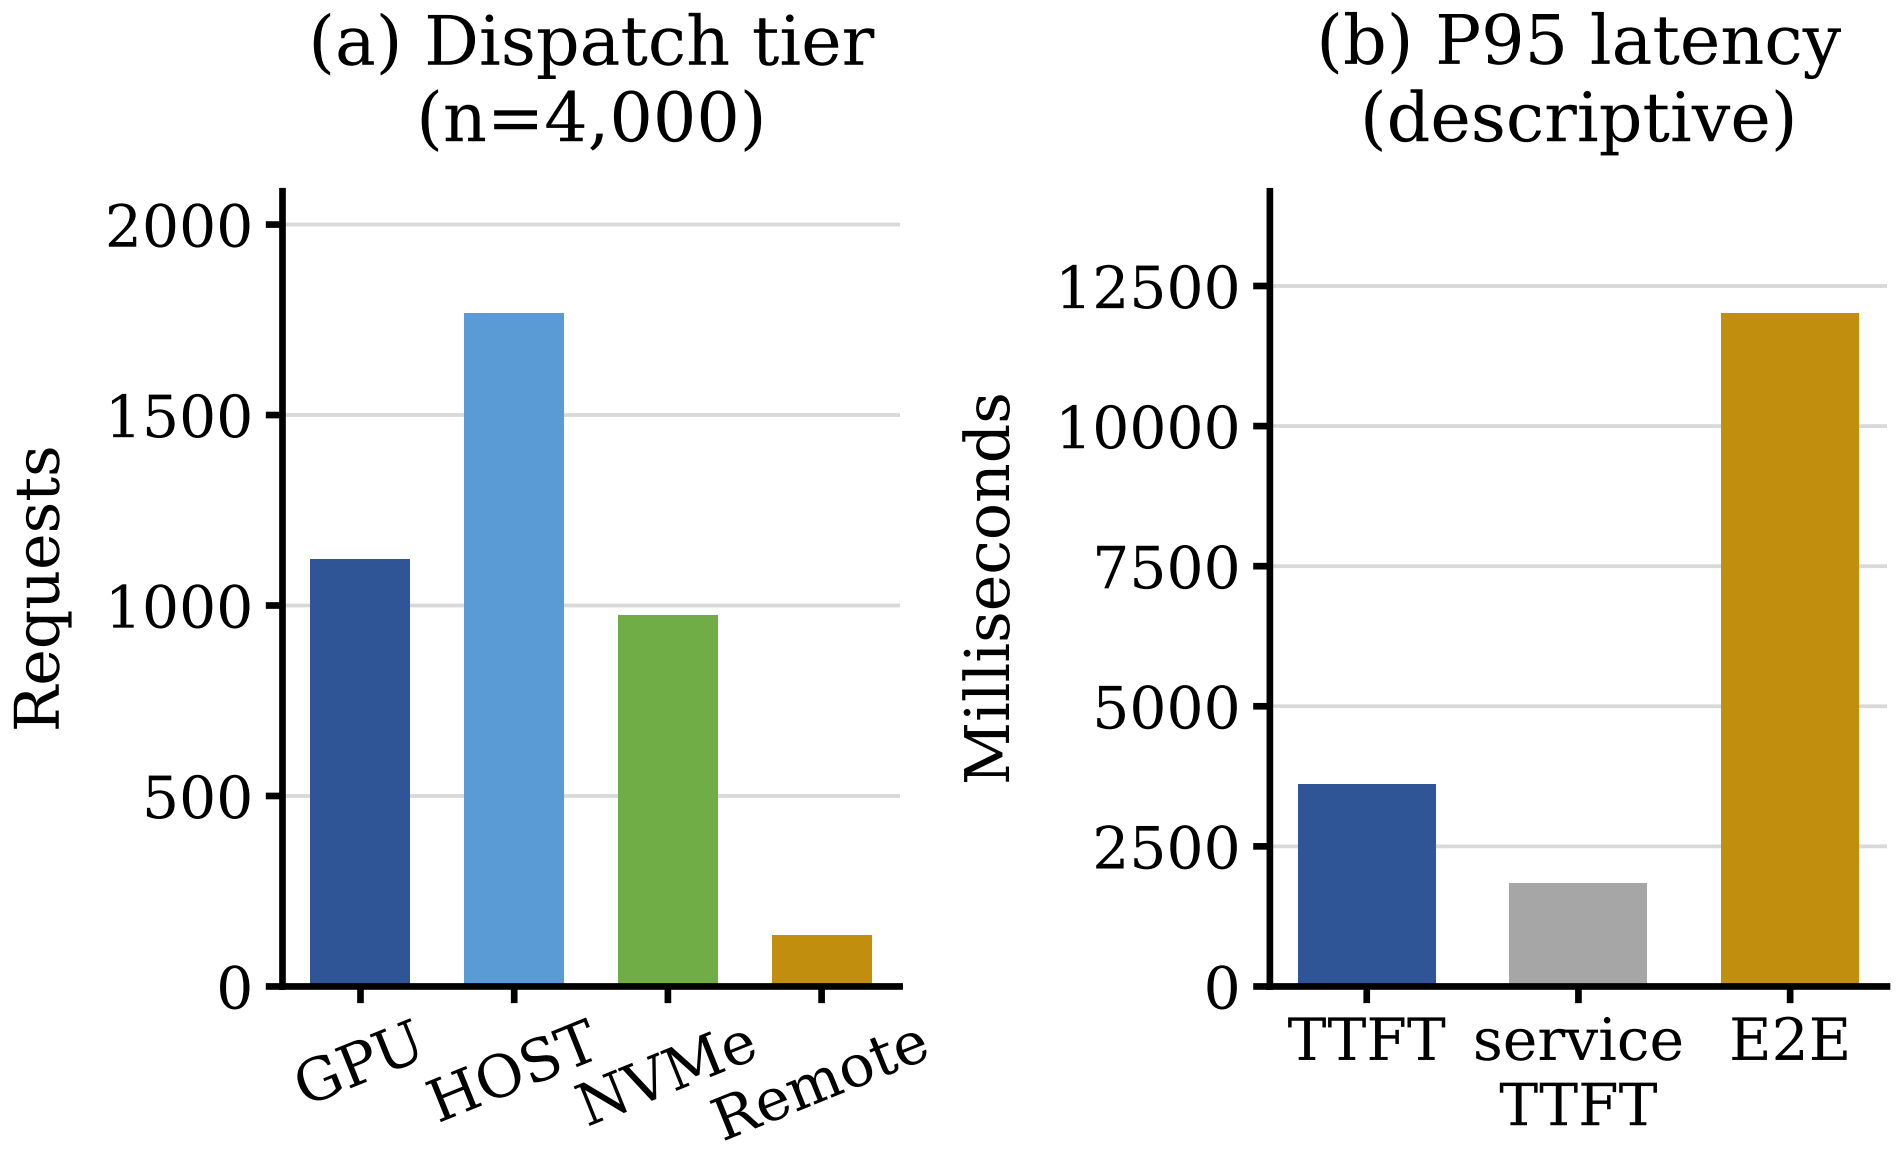
(a) Dispatch tier
(n=4,000)
Requests
2000
1500
1000
500
0
GPU
HOST
NVMe
Remote
(b) P95 latency
(descriptive)
Milliseconds
12500
10000
7500
5000
2500
0
TTFT
service
TTFT
E2E